

Distributed Systems in Action

Mohamed Sweelam

Software Engineer

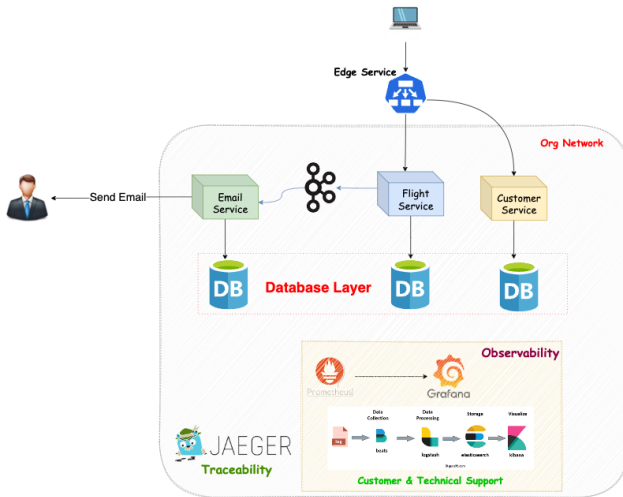
Outline

- 1 Introduction to Distributed Systems
- 2 Use Case: Flight System
- 3 CAP Theorem
- 4 Distributed Databases
- 5 High Availability and Scalability
- 6 Scalability
- 7 Resiliency
- 8 Data Replication
- 9 Distributed Messaging Systems
- 10 Distributed Key-Value Stores
- 11 Consensus
- 12 Physical Clock & Logical Clock
- 13 Distributed System Security
- 14 Distributed System Monitoring
- 15 Conclusion

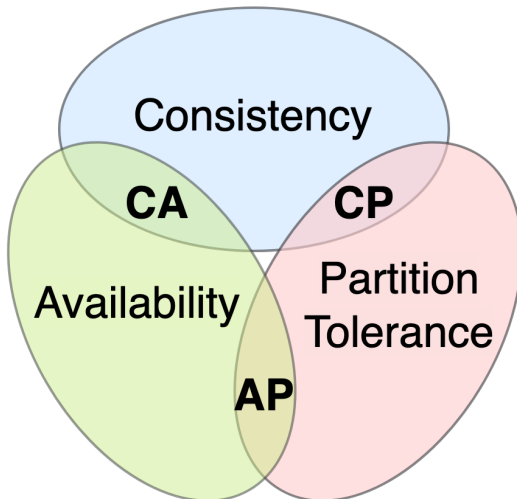
Introduction to Distributed Systems

- Definition and characteristics of distributed systems
- Importance and benefits of distributed systems
- Challenges and trade-offs in designing distributed systems

Use Case: Flight System

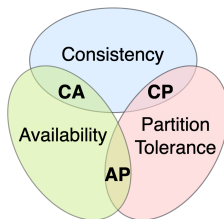


CAP Theorem



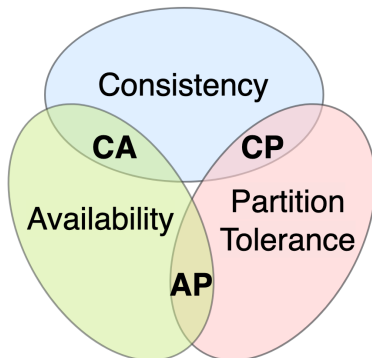
Distributed Databases

- **Consistency** every read operation gives a result of the most recent write.
- **Availability** every read operation gives a non-error response, but, data might be staled.
- **Partition Tolerance** system operates normally despite network failure.



Distributed Databases

In the presence of network partitioning, system designers must choose between data consistency and availability.



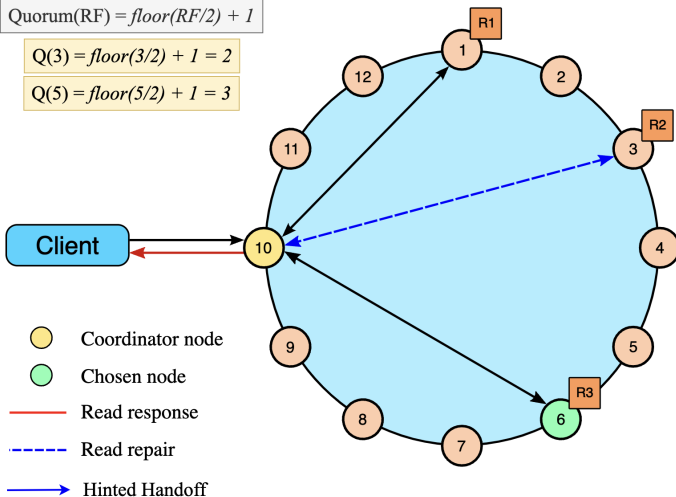
Distributed Databases

Coordinator waits until Quorum acknowledge

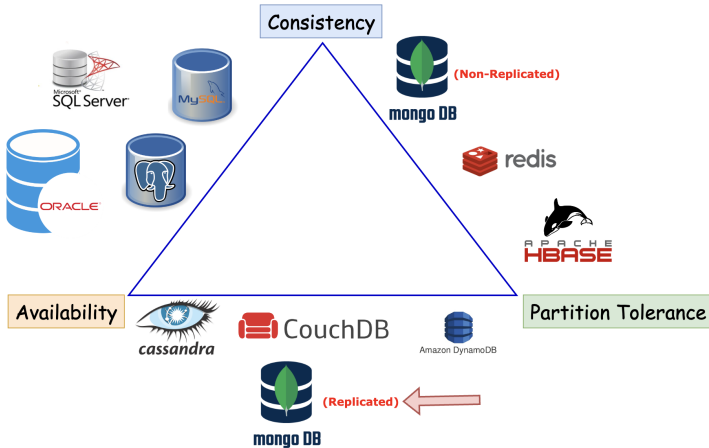
$$\text{Quorum}(\text{RF}) = \text{floor}(\text{RF}/2) + 1$$

$$Q(3) = \text{floor}(3/2) + 1 = 2$$

$$Q(5) = \text{floor}(5/2) + 1 = 3$$

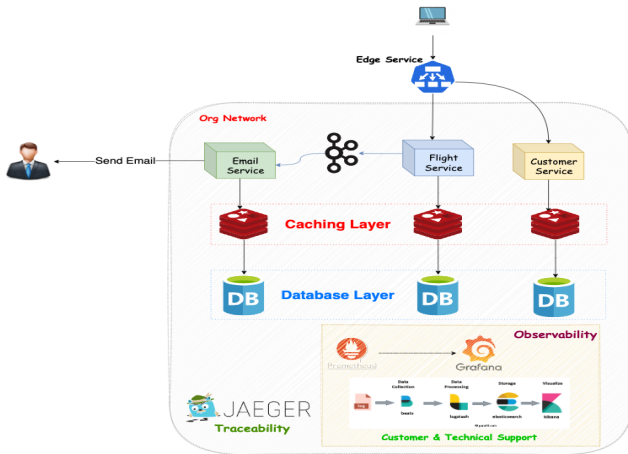


Distributed Databases



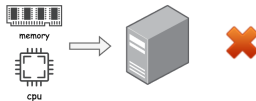
High Availability and Scalability

- Load Balancer
- API Gateway
- Rate Limiter

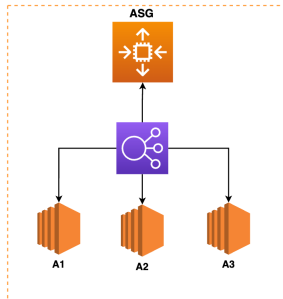
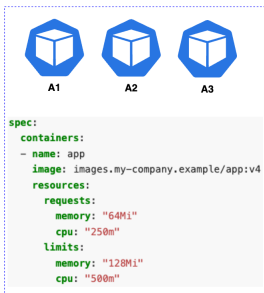


Horizontal Scaling vs Vertical Scaling

Vertical Scaling

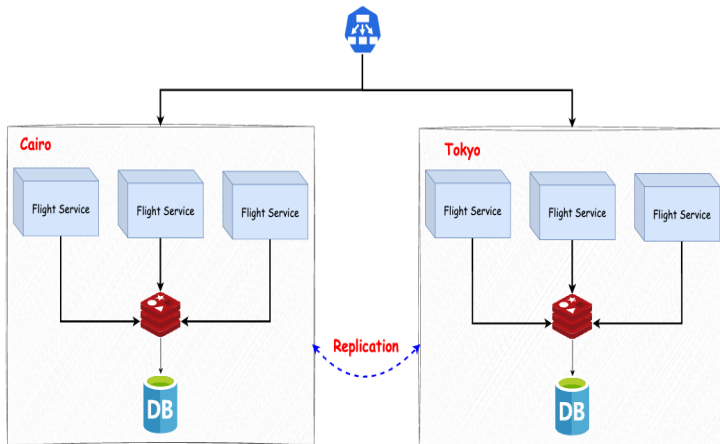


Horizontal Scaling



```
spring:
  cloud:
    default-filters:
      - name: Retry
        args:
          retry: 5
          methods: GET
          backoff:
            firstBackOff: 10ms
            maxBackOff: 50ms
            factor: 2
            basedOnPreviousValue: false
```

Data Replication


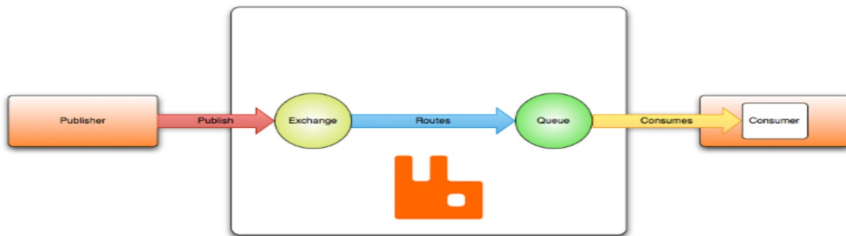


Distributed Messaging Systems

- Introduction to distributed messaging systems
- Overview of Kafka and RabbitMQ

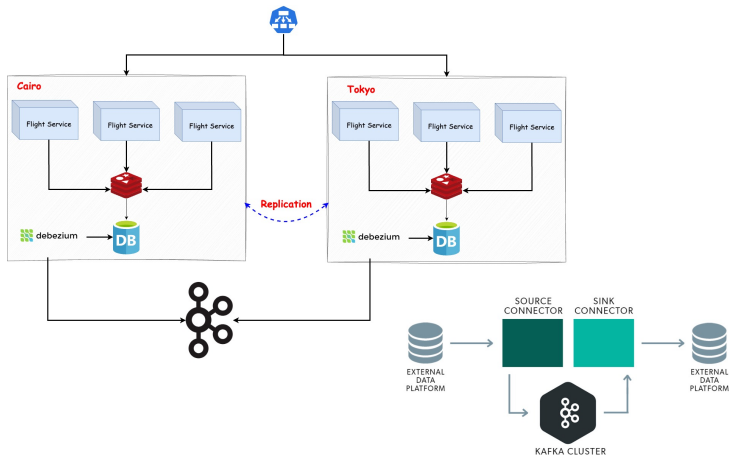
RabbitMQ vs Kafka

"Hello, world" example routing



Data Replication: CDC

In databases, **change data capture (CDC)** is a set of *software design patterns* used to determine and track the data that has changed (the "deltas") so that action can be taken using the changed data.



Data Replication: CDC

```
{
  "schema": {},
  "payload": {
    "source": {
      "version": "2.2.1.Final",
      "connector": "mysql",
      "name": "mysql",
      "db": "inventory",
      "sequence": null,
      "table": "customers"
    },
    "databaseName": "inventory",
    "schemaName": null,
    "ddl": "ALTER TABLE customers ADD middle_name varchar(255) AFTER first_name",
    "tableChanges": [
      {
        "type": "ALTER",
        "id": "\\\"inventory\\\"\\.\\\"customers\\\"\\.\\\"",
        "table": {
          "defaultCharsetName": "utf8mb4",
          "primaryKeyColumnNames": [
            "id"
          ]
        },
        "columns": [
          {
            "name": "id",
            "jdbcType": 4,
            "nativeType": null,
            "typeName": "INT",
            "typeExpression": "INT",
            "charsetName": null,
            "length": null,
            "scale": null,
            "position": 1,
            "optional": false,
            "autoIncremented": true,
            "generated": true
          }
        ]
      }
    ]
  }
}
```

Distributed Messaging Systems

- Introduction to distributed messaging systems
- Overview of Kafka and RabbitMQ
- Comparison between Kafka and RabbitMQ:

| Kafka | RabbitMQ |
|--|--|
| High-throughput, fault-tolerant distributed streaming platform. | Robust and flexible messaging broker. |
| Emphasizes real-time event streaming and data pipeline use cases. | Implements Advanced Message Queuing Protocol (AMQP). |
| Provides strong durability and replication guarantees. | Focuses on message queuing and asynchronous communication. |
| Scales horizontally to handle large-scale data streams. | Provides various messaging patterns (e.g., publish-subscribe, point-to-point). |
| Supports complex event processing with built-in stream processing. | Offers pluggable message durability, routing, and acknowledgement mechanisms. |

Distributed Key-Value Stores

A type of NoSQL DBs, the idea behind having **it** is having an extremely fast mechanism for retrieving the data.

- Introduction to distributed key-value stores (e.g., Redis, Memcached)
- Data partitioning and replication techniques
- Consistency and availability trade-offs

Distributed Key-Value Stores: Memcached vs Redis

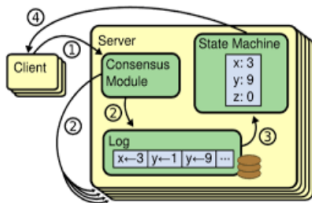


| Memcached | Redis |
|--|---|
| In-memory key-value store primarily designed for caching. | In-memory key-value store with additional data structures and functionality. |
| Supports simple key-value operations like GET, SET, DELETE. | Offers a rich set of data structures (strings, lists, sets, hashes, sorted sets). |
| Lacks built-in persistence mechanisms. | Provides persistence options (snapshotting, append-only file, replication). |
| Does not support complex queries or secondary indexes. | Supports more advanced operations like sorting, ranking, and filtering. |
| Focuses on high performance and low latency. | Balances performance with additional functionality and data structures. |
| No built-in way of replication, requires additional component. | Come with out-of-the-box replication using cluster, and as feature with Sentinel . |

Consensus

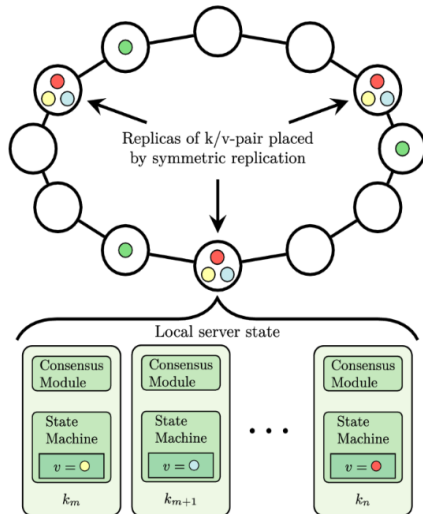
Replicated State Machine

Implementing a fault-tolerant service by replicating servers and coordinating client interactions with server replicas



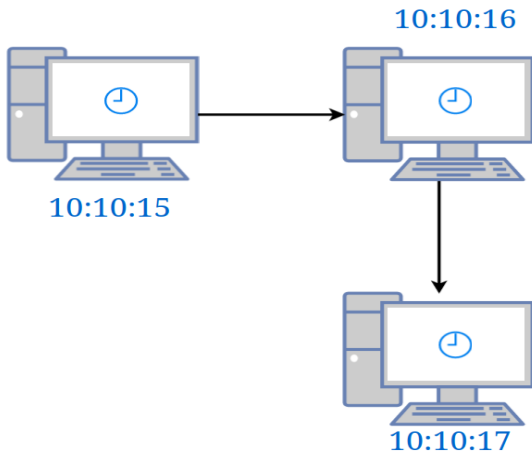
Algorithms

- **Paxos**
- **Raft**




Natural/Physical Clock

What is the order of the events?



Natural/Physical Clock

What is the order of the events?



Natural/Physical Clock

Quartz Clock



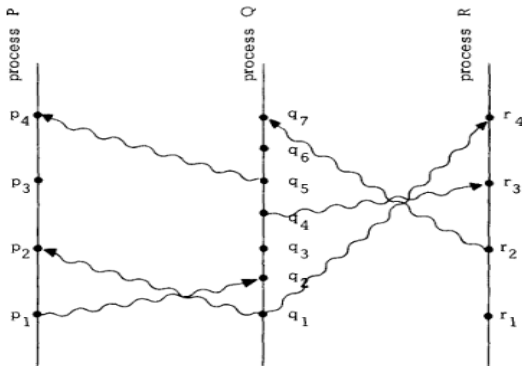
Logical Clock (Lamport Algorithm)

- Provide **happened before** relationship

$$a \rightarrow b$$

- Total Ordering

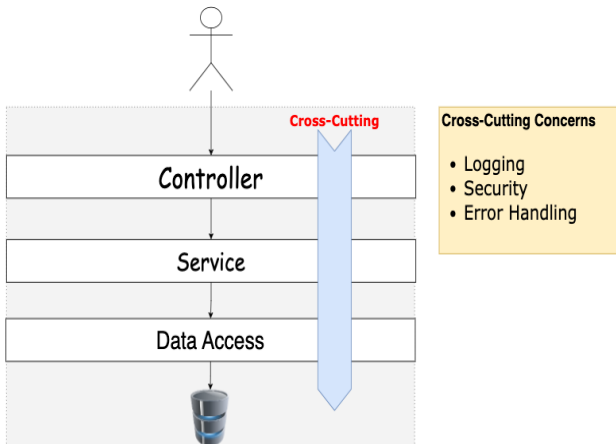
$$a \Rightarrow b \iff a \rightarrow b$$



Distributed System Security



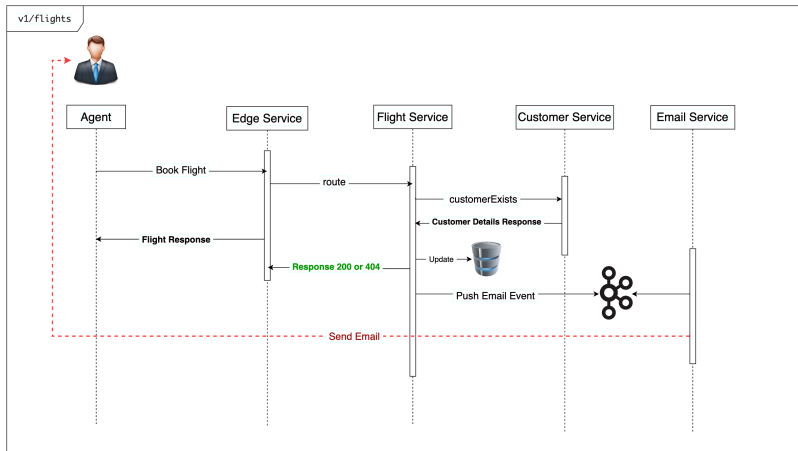
Distributed System Security



Distributed System Monitoring

- Log aggregation and distributed tracing
- Metrics collection and monitoring tools (e.g., Prometheus, Grafana)
- Anomaly detection and performance optimization

API Documentations



Conclusion

- Recap of distributed systems concepts
- Overview of various distributed systems topics
- Further resources for exploring distributed systems in depth